

Group Project 9660

Group 2

05-10-2023

Install packages

```
#install.packages("dplyr")
#install.packages("stargazer")
#install.packages("rmarkdown")
#install.packages("stargazer")
#install.packages("car")
#install.packages("ggplot2")
#install.packages("rlang", dependencies = TRUE)
#install.packages("tinytex")
#tinytex::install_tinytex()
#install.packages("stats")
#install.packages("psych")
#install.packages("GGally")
#install.packages("tidyverse")
#install.packages("tidymodels")
#install.packages("forcats")
#install.packages("corrplot")
#install.packages("mgcv")
```

Original Data preview

```
# Load the dataset
Original_Video.Game <- read.csv("Video_Games_Sales_as_at_22_Dec_2016.csv")

# Filter the data for years between 2000 and 2016
Original_Video.Game <- Original_Video.Game %>%
filter(Year_of_Release >= 2000 & Year_of_Release <= 2016)

# Print the summary of the dataset
summary(Original_Video.Game)
```

##	Name	Platform	Year_of_Release	Genre
Publisher	NA_Sales	EU_Sales	JP_Sales	
##	Length:14470	Length:14470	Length:14470	Length:14470
Length:14470	Min. : 0.0000	Min. : 0.0000	Min. : 0.00000	
##	Class :character	Class :character	Class :character	Class :character
Class :character	1st Qu.: 0.0000	1st Qu.: 0.0000	1st Qu.: 0.00000	
##	Mode :character	Mode :character	Mode :character	Mode :character
Mode :character	Median : 0.0800	Median : 0.0200	Median : 0.00000	
##				

```

Mean : 0.2439 Mean : 0.1441 Mean :0.05638
##
3rd Qu.: 0.2300 3rd Qu.: 0.1100 3rd Qu.:0.03000
##
Max. :41.3600 Max. :28.9600 Max. :6.50000
##
## Other_Sales Global_Sales Critic_Score Critic_Count
User_Score User_Count Developer Rating
## Min. : 0.00000 Min. : 0.010 Min. :13.00 Min. : 3.0
Length:14470 Min. : 4.0 Length:14470 Length:14470
## 1st Qu.: 0.00000 1st Qu.: 0.060 1st Qu.:60.00 1st Qu.: 12.0 Class
:character 1st Qu.: 10.0 Class :character Class :character
## Median : 0.01000 Median : 0.160 Median :71.00 Median : 22.0 Mode
:character Median : 24.0 Mode :character Mode :character
## Mean : 0.05032 Mean : 0.495 Mean :68.86 Mean : 26.6
Mean : 161.5
## 3rd Qu.: 0.04000 3rd Qu.: 0.440 3rd Qu.:79.00 3rd Qu.: 37.0
3rd Qu.: 81.0
## Max. :10.57000 Max. :82.530 Max. :98.00 Max. :113.0
Max. :10665.0
## NA's :6583 NA's :6583
NA's :7099

# Count the number of missing values in each column
original_missing_values <- summarise_all(Original_Video.Game, list(~
sum(is.na(.))))
print(original_missing_values)
## Name Platform Year_of_Release Genre Publisher NA_Sales EU_Sales JP_Sales
Other_Sales Global_Sales Critic_Score Critic_Count User_Score User_Count
## 1 0 0 0 0 0 0 0 0 0
0 0 6583 6583 0 7099
## Developer Rating
## 1 0 0

# Print the dimensions of the dataset
dim(Original_Video.Game)
## [1] 14470 16

# View the dataset in a separate window
View(Original_Video.Game)

# Apply a function to return the class of each column
sapply(Original_Video.Game,class)
## Name Platform Year_of_Release Genre
Publisher NA_Sales EU_Sales JP_Sales Other_Sales
## "character" "character" "character" "character"
"character" "numeric" "numeric" "numeric" "numeric"
## Global_Sales Critic_Score Critic_Count User_Score
User_Count Developer Rating

```

```
##           "numeric"           "integer"           "integer"           "character"
"integer"           "character"           "character"

# Print the structure of the dataset
str(Original_Video.Game)
## 'data.frame':   14470 obs. of  16 variables:
##  $ Name           : chr  "Wii Sports" "Mario Kart Wii" "Wii Sports Resort"
##  "New Super Mario Bros." ...
##  $ Platform       : chr  "Wii" "Wii" "Wii" "DS" ...
##  $ Year_of_Release: chr  "2006" "2008" "2009" "2006" ...
##  $ Genre          : chr  "Sports" "Racing" "Sports" "Platform" ...
##  $ Publisher      : chr  "Nintendo" "Nintendo" "Nintendo" "Nintendo" ...
##  $ NA_Sales       : num  41.4 15.7 15.6 11.3 14 ...
##  $ EU_Sales       : num  28.96 12.76 10.93 9.14 9.18 ...
##  $ JP_Sales       : num  3.77 3.79 3.28 6.5 2.93 4.7 1.93 4.13 3.6 0.24
##  ...
##  $ Other_Sales    : num  8.45 3.29 2.95 2.88 2.84 2.24 2.74 1.9 2.15 1.69
##  ...
##  $ Global_Sales   : num  82.5 35.5 32.8 29.8 28.9 ...
##  $ Critic_Score   : int  76 82 80 89 58 87 NA 91 80 61 ...
##  $ Critic_Count   : int  51 73 73 65 41 80 NA 64 63 45 ...
##  $ User_Score     : chr  "8" "8.3" "8" "8.5" ...
##  $ User_Count     : int  322 709 192 431 129 594 NA 464 146 106 ...
##  $ Developer      : chr  "Nintendo" "Nintendo" "Nintendo" "Nintendo" ...
##  $ Rating         : chr  "E" "E" "E" "E" ...

# reference the original variables in the data frame
attach(Original_Video.Game)
```

Original Data statistics

```
# Remove outliers using the IQR method for Global_Sales
q1 <- quantile(Original_Video.Game$Global_Sales, 0.25, na.rm = TRUE)
q3 <- quantile(Original_Video.Game$Global_Sales, 0.75, na.rm = TRUE)
iqr <- q3 - q1
upper <- q3 + 1.5*iqr
lower <- q1 - 1.5*iqr
Original_Video.Game <- subset(Original_Video.Game, Global_Sales >= lower &
Global_Sales <= upper)

# calculate mean for global sales
original_global_sales_mean <- mean(Original_Video.Game$Global_Sales)
print(original_global_sales_mean)
## [1] 0.2142826

# Calculate variance of a numeric variable
original_variance <- var(Original_Video.Game$Global_Sales, na.rm = TRUE)
print(original_variance)
## [1] 0.05088653
```

```


# Calculate standard deviation of a numeric variable
original_sd <- sd(Original_Video.Game$Global_Sales, na.rm = TRUE)
print(original_sd)
## [1] 0.2255804

# Set the seed for reproducibility
set.seed(3)

# Add random normal values to the Global_Sales variable
Original_Video.Game$Global_Sales <- Original_Video.Game$Global_Sales +
rnorm(nrow(Original_Video.Game), mean = original_global_sales_mean , sd =
original_sd)

# Plot a histogram of the Global_Sales variable
ggplot(Original_Video.Game, aes(x = Global_Sales)) +
  geom_histogram(bins = 20, color = "black", fill = "lightblue") +
  labs(title = "Global Sales Distribution", x = "Global Sales", y =
"Frequency")

```



```

# Calculate correlation between numeric variables
original_correlations <- cor(Original_Video.Game[,
sapply(Original_Video.Game, is.numeric)], use = "complete.obs")
print(original_correlations)
##           NA_Sales  EU_Sales  JP_Sales  Other_Sales

```

```

Global_Sales Critic_Score Critic_Count User_Count
## NA_Sales      1.0000000000 0.3379415 0.0005509181 0.49594706
0.61576357 0.17868864 0.2157053 -0.02426960
## EU_Sales      0.3379415234 1.0000000 0.0128536964 0.63192318
0.54957274 0.20083235 0.2654758 0.27087742
## JP_Sales      0.0005509181 0.0128537 1.0000000000 0.04906703
0.21069477 0.08123742 0.1184602 -0.01777846
## Other_Sales   0.4959470639 0.6319232 0.0490670289 1.00000000
0.53433409 0.13560410 0.2259492 0.09626510
## Global_Sales  0.6157635723 0.5495727 0.2106947663 0.53433409
1.00000000 0.16508905 0.2379256 0.08878281
## Critic_Score  0.1786886365 0.2008323 0.0812374237 0.13560410
0.16508905 1.00000000 0.3306252 0.21331815
## Critic_Count  0.2157052746 0.2654758 0.1184601987 0.22594917
0.23792564 0.33062523 1.0000000 0.22654774
## User_Count    -0.0242695983 0.2708774 -0.0177784646 0.09626510
0.08878281 0.21331815 0.2265477 1.00000000

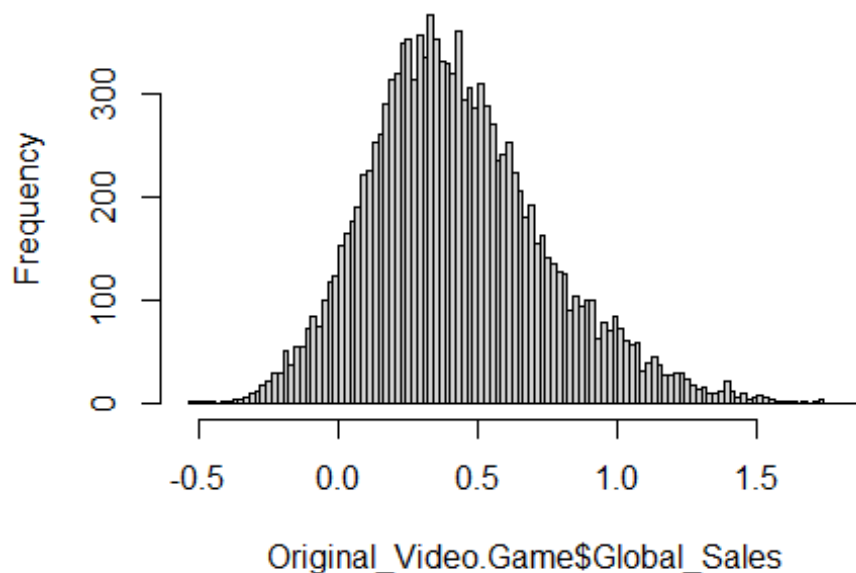
```

```

# Plot histogram of a numeric variable
hist(Original_Video.Game$Global_Sales, breaks = 100)

```

Histogram of Original_Video.Game\$Global_Sales




```

# Plot a density plot of Global_Sales
plot(density(Original_Video.Game$Global_Sales))

```


```
density.default(x = Original_Video.Game$Global_Sa
```



```
# Calculate the mean for numeric columns
sapply(Original_Video.Game[sapply(Original_Video.Game, is.numeric)], mean)
##      NA_Sales      EU_Sales      JP_Sales      Other_Sales      Global_Sales
Critic_Score Critic_Count  User_Count
##  0.11042469  0.05381988  0.03089674  0.01882531  0.42550297
NA           NA           NA

# Calculate the standard deviation for numeric columns
sapply(Original_Video.Game[sapply(Original_Video.Game, is.numeric)], sd)
##      NA_Sales      EU_Sales      JP_Sales      Other_Sales      Global_Sales
Critic_Score Critic_Count  User_Count
##  0.14206790  0.09111080  0.08295832  0.03453205  0.31994910
NA           NA           NA

# Create a density plot of Global_Sales
ggplot(Original_Video.Game, aes(Global_Sales)) +
  geom_density(fill = "#69b3a2", alpha = 0.8) +
  ggtitle("Density Plot of Global Sales") +
  xlab("Global Sales") +
  ylab("Density")
```



Revised Data Preview

Remove missing values

```
Video.Game <- na.omit(Original_Video.Game)
```

Convert categorical variables to factors

```
Video.Game$Name <- as.factor(Video.Game$Name)
```

```
Video.Game$Platform <- as.factor(Video.Game$Platform)
```

```
Video.Game$Genre <- as.factor(Video.Game$Genre)
```

```
Video.Game$Publisher <- as.factor(Video.Game$Publisher)
```

```
Video.Game$Developer <- as.factor(Video.Game$Developer)
```

```
Video.Game$Rating <- as.factor(Video.Game$Rating)
```

Convert numeric variables to numeric

```
Video.Game$Year_of_Release <- as.numeric(Video.Game$Year_of_Release)
```

```
Video.Game$Critic_Score <- as.numeric(Video.Game$Critic_Score)
```

```
Video.Game$Critic_Count <- as.numeric(Video.Game$Critic_Count)
```

```
Video.Game$User_Score <- as.numeric(Video.Game$User_Score)
```

```
Video.Game$User_Count <- as.numeric(Video.Game$User_Count)
```

Filter the data for years between 2006 and 2016

```
Video.Game <- Video.Game %>%
```

```
filter(Year_of_Release >= 2000 & Year_of_Release <= 2016)
```

Print the summary of the dataset

```
summary(Video.Game)
```

```
##                               Name           Platform
Year_of_Release      Genre      Publisher
## Harry Potter and the Order of the Phoenix:  7   PS2       : 880   Min.
:2000   Action      :1341   Electronic Arts : 673
## Spider-Man 3                               :   7   X360    : 651   1st
Qu.:2004   Sports      : 755   Ubisoft      : 412
## Terraria                               :   7   PC       : 633   Median
:2007   Shooter      : 663   Activision    : 365
## Tomb Raider: Legend                       :   7   PS3     : 564   Mean
:2008   Role-Playing: 581   THQ           : 263
## FIFA World Cup Germany 2006                :   6   XB      : 519   3rd
Qu.:2011   Racing      : 467   Sega         : 236
## Ghostbusters: The Video Game                :   6   DS      : 385   Max.
:2016   Platform      : 310   Sony Computer Entertainment: 221
## (Other)                               :5512   (Other):1920
(Other)      :1435   (Other)                :3382
```

```
##      NA_Sales      EU_Sales      JP_Sales      Other_Sales
Global_Sales      Critic_Score      Critic_Count      User_Score
## Min.      :0.0000   Min.      :0.0000   Min.      :0.0000   Min.      :0.00000
Min.      :-0.4531   Min.      :13.00   Min.      : 3.00   Min.      :0.500
## 1st Qu.:0.0400   1st Qu.:0.01000   1st Qu.:0.00000   1st Qu.:0.01000
1st Qu.: 0.2546   1st Qu.:60.00   1st Qu.: 13.00   1st Qu.:6.300
## Median :0.1100   Median :0.04000   Median :0.00000   Median :0.01500
Median : 0.4704   Median :70.00   Median : 22.00   Median :7.400
## Mean      :0.1582   Mean      :0.08054   Mean      :0.02185   Mean      :0.02789
Mean      : 0.5009   Mean      :68.03   Mean      : 25.64   Mean      :7.083
## 3rd Qu.:0.2200   3rd Qu.:0.11000   3rd Qu.:0.00000   3rd Qu.:0.04000
3rd Qu.: 0.7143   3rd Qu.:78.00   3rd Qu.: 35.00   3rd Qu.:8.200
## Max.      :0.9400   Max.      :0.93000   Max.      :0.74000   Max.      :0.54000
Max.      : 1.8782   Max.      :96.00   Max.      :106.00   Max.      :9.600
```

```
##
##      User_Count      Developer      Rating
## Min.      :    4.0   EA Canada : 108      : 64
## 1st Qu.:    9.0   EA Sports  : 96      E :1616
## Median :   21.0   Capcom    : 91     E10+ :784
## Mean      :  104.1   Ubisoft   : 83      M :1091
## 3rd Qu.:   57.0   Konami    : 80      RP : 1
## Max.      :10665.0   Omega Force: 64      T :1996
##
##      (Other)      :5030
```

```
# Count the number of missing values in each column
```

```
revised_missing_values <- summarise_all(Video.Game, list(~ sum(is.na(.))))
print(revised_missing_values)
```

```
##      Name Platform Year_of_Release Genre Publisher NA_Sales EU_Sales JP_Sales
Other_Sales Global_Sales Critic_Score Critic_Count User_Score User_Count
## 1      0      0      0      0      0      0      0      0      0
0      0      0      0      0      0      0
##      Developer Rating
## 1      0      0
```



```

# Print the dimensions of the dataset
dim(Video.Game)
## [1] 5552 16

# View the dataset in a separate window
View(Video.Game)

# Apply a function to return the class of each column
sapply(Video.Game,class)
##           Name           Platform Year_of_Release           Genre
Publisher      NA_Sales      EU_Sales      JP_Sales      Other_Sales
##           "factor"           "factor"           "numeric"           "factor"
"factor"           "numeric"           "numeric"           "numeric"           "numeric"
##   Global_Sales Critic_Score Critic_Count      User_Score
User_Count      Developer      Rating
##           "numeric"           "numeric"           "numeric"           "numeric"
"numeric"           "factor"           "factor"

# Print the structure of the dataset
str(Video.Game)
## 'data.frame': 5552 obs. of 16 variables:
## $ Name : Factor w/ 3793 levels " Tales of Xillia 2",...: 3599
2064 34 3296 2254 185 2531 2432 3780 2408 ...
## $ Platform : Factor w/ 17 levels "3DS","DC","DS",...: 13 13 15 13 15
15 10 15 14 14 ...
## $ Year_of_Release: num 2008 2007 2007 2007 2008 ...
## $ Genre : Factor w/ 12 levels "Action","Adventure",...: 5 3 10 1
1 1 5 1 1 3 ...
## $ Publisher : Factor w/ 256 levels "10TACLE Studios",...: 229 231 23
64 224 236 236 236 236 155 ...
## $ NA_Sales : num 0.51 0.44 0.69 0.45 0.65 0.58 0.22 0.54 0.52 0.5
...
## $ EU_Sales : num 0.4 0.46 0.04 0.46 0.22 0.34 0.64 0.34 0.36 0.26
...
## $ JP_Sales : num 0 0 0.22 0 0.05 0 0 0.02 0.05 0.17 ...
## $ Other_Sales : num 0.11 0.11 0.06 0.11 0.1 0.09 0.16 0.1 0.08 0.08
...
## $ Global_Sales : num 1.007 1.158 0.964 1.268 1.476 ...
## $ Critic_Score : num 51 74 80 64 81 72 90 81 77 76 ...
## $ Critic_Count : num 18 23 54 19 74 32 16 73 70 74 ...
## $ User_Score : num 3.6 8 7.9 6.1 8 7.3 8.5 7 7.7 8 ...
## $ User_Count : num 8 27 101 38 302 318 525 189 758 273 ...
## $ Developer : Factor w/ 1224 levels "", "10tacle Studios,
Fusionsphere Systems",...: 481 28 832 356 1066 1128 1135 1136 1135 115 ...
## $ Rating : Factor w/ 6 levels "", "E", "E10+",...: 2 6 6 6 4 4 3 6 4
3 ...
## - attr(*, "na.action")= 'omit' Named int [1:7328] 3 6 7 11 18 19 20 22 23
24 ...

```

```
## ... attr(*, "names")= chr [1:7328] "1593" "1596" "1597" "1601" ...

# reference the revised variables in the data frame
attach(Video.Game)
## The following objects are masked from Original_Video.Game:
##
## Critic_Count, Critic_Score, Developer, EU_Sales, Genre, Global_Sales,
## JP_Sales, NA_Sales, Name, Other_Sales, Platform, Publisher,
## Rating, User_Count, User_Score, Year_of_Release
```

Revised Data statistics

```
# Remove outliers using the IQR method for Global_Sales
q1 <- quantile(Video.Game$Global_Sales, 0.25, na.rm = TRUE)
q3 <- quantile(Video.Game$Global_Sales, 0.75, na.rm = TRUE)
iqr <- q3 - q1
upper <- q3 + 1.5*iqr
lower <- q1 - 1.5*iqr
Video.Game <- subset(Video.Game, Global_Sales >= lower & Global_Sales <=
upper)

# calculate mean for global sales
revised_global_sales_mean <- mean(Video.Game$Global_Sales)
print(revised_global_sales_mean)
## [1] 0.4935089

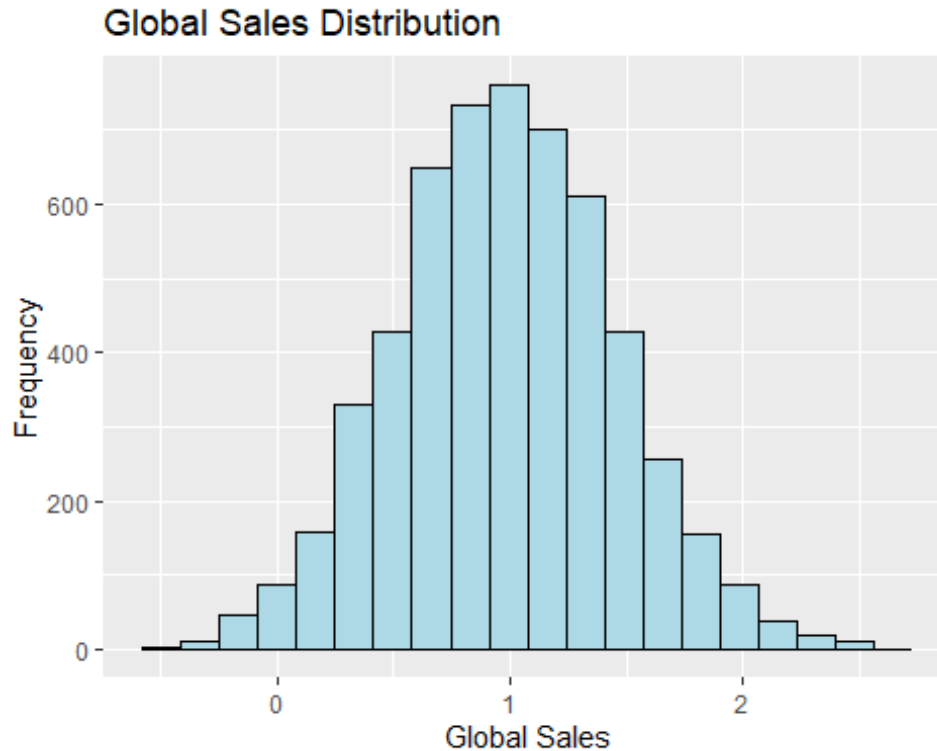
# Calculate variance of a numeric variable
revised_variance <- var(Video.Game$Global_Sales, na.rm = TRUE)
print(revised_variance)
## [1] 0.1088451

# Calculate standard deviation of a numeric variable
revised_sd <- sd(Video.Game$Global_Sales, na.rm = TRUE)
print(revised_sd)
## [1] 0.3299168

# Set the seed for reproducibility
set.seed(3)

# Add random normal values to the Global_Sales variable
Video.Game$Global_Sales <- Video.Game$Global_Sales + rnorm(nrow(Video.Game),
mean = revised_global_sales_mean , sd = revised_sd)

# Plot a histogram of the Global_Sales variable
ggplot(Video.Game, aes(x = Global_Sales)) +
  geom_histogram(bins = 20, color = "black", fill = "lightblue") +
  labs(title = "Global Sales Distribution", x = "Global Sales", y =
"Frequency")
```




```
# Calculate correlation between numeric variables
revised_correlations <- cor(Video.Game[, sapply(Video.Game, is.numeric)], use
= "complete.obs")
print(revised_correlations)
```

	Year_of_Release	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score	Critic_Count	User_Score	User_Count
Year_of_Release	1.000000000	-0.101666790	0.09731540	0.050232518	0.10942815	-0.01339470	0.002958213	0.1356116	-0.22342259	0.18560796
NA_Sales	-0.101666790	1.000000000	0.33073898	0.003311067	0.49309080	0.43473854	0.172473322	0.2119618	0.07768621	-0.03019462
EU_Sales	0.097315396	0.330738982	1.000000000	0.011276202	0.62099581	0.37426514	0.194451819	0.2646900	0.03791376	0.26882240
JP_Sales	0.050232518	0.003311067	0.01127620	1.000000000	0.04698147	0.15200756	0.080851048	0.1195634	0.13178419	-0.01729629
Other_Sales	0.109428152	0.493090798	0.62099581	0.046981468	1.000000000	0.37918641	0.128752641	0.2231947	0.03404031	0.08994499
Global_Sales	-0.013394702	0.434738538	0.37426514	0.152007564	0.37918641	1.000000000	0.103088284	0.1752382	0.04843587	0.07126432
Critic_Score	0.002958213	0.172473322	0.19445182	0.080851048	0.12875264	0.10308828	1.000000000	0.3292032	0.58866427	0.21316760
Critic_Count	0.135611580	0.211961787	0.26468996	0.119563431	0.22319474	0.17523816	0.329203202	1.00000000	0.18925007	0.23206714
User_Score	-0.223422592	0.077686207	0.03791376	0.131784191	0.03404031	0.04843587	0.588664269	0.1892501	1.00000000	0.02890343
User_Count	0.185607961	-0.030194616	0.26882240	-0.017296294						


```
0.08994499 0.07126432 0.213167603 0.2320671 0.02890343 1.00000000
```

```
# Plot histogram of a numeric variable  
hist(Video.Game$Global_Sales, breaks = 100)
```



```
# Plot a density plot of Global_Sales  
plot(density(Video.Game$Global_Sales))
```

density.default(x = Video.Game\$Global_Sales)




N = 5510 Bandwidth = 0.0754

```
# Calculate the mean for numeric columns
sapply(Video.Game[sapply(Video.Game, is.numeric)], mean)
## Year_of_Release      NA_Sales      EU_Sales      JP_Sales
Other_Sales Global_Sales Critic_Score Critic_Count User_Score
## 2.007553e+03 1.559837e-01 7.898004e-02 2.173684e-02
2.738475e-02 9.812912e-01 6.796225e+01 2.557423e+01 7.081579e+00
## User_Count
## 1.027194e+02

# Calculate the standard deviation for numeric columns
sapply(Video.Game[sapply(Video.Game, is.numeric)], sd)
## Year_of_Release      NA_Sales      EU_Sales      JP_Sales
Other_Sales Global_Sales Critic_Score Critic_Count User_Score
## 4.11060136 0.15507885 0.10081054 0.06792289
0.03595749 0.46918224 13.57073499 16.22036452 1.46796618
## User_Count
## 388.33043602

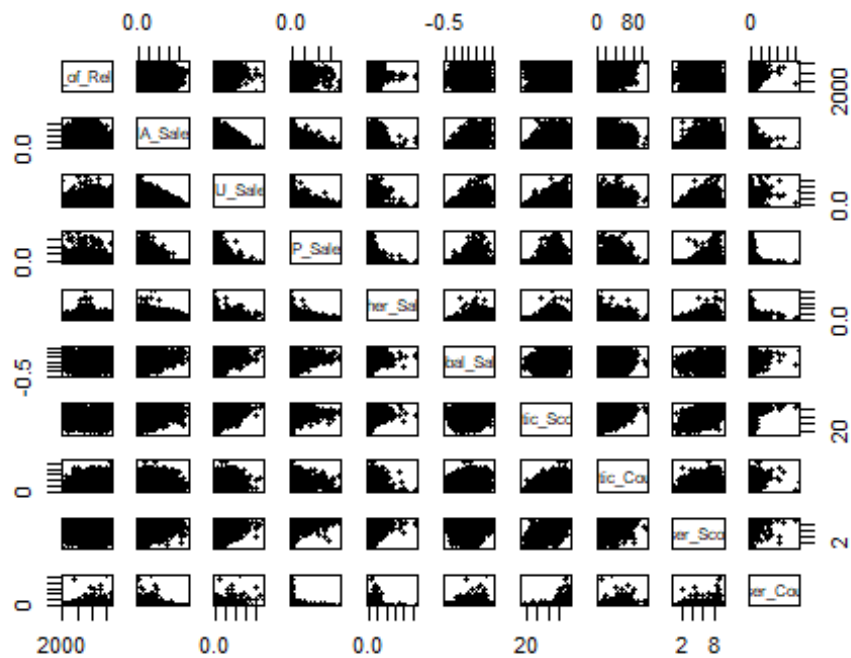
# Create a density plot of Global_Sales
ggplot(Video.Game, aes(Global_Sales)) +
  geom_density(fill = "#69b3a2", alpha = 0.8) +
  ggtitle("Density Plot of Global Sales") +
  xlab("Global Sales") +
  ylab("Density")
```



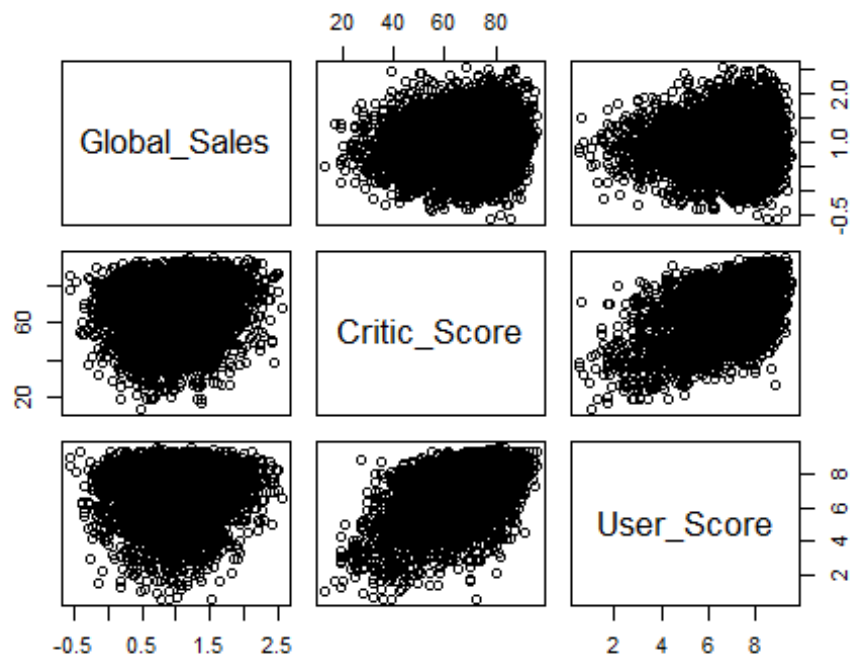
Revised Data

statistics


```
# Select only the numeric columns  
num_data <- Video.Game %>% select_if(is.numeric)  
  
# Create scatterplot matrix using pairs() function  
pairs(num_data, pch = 20)
```




Create scatterplot matrix using subset of variables and pairs() function
 pairs(~Global_Sales + Critic_Score + User_Score, Video.Game)



```
# Plot using a subset of variables  
plot(Critic_Score,Global_Sales)
```



```
plot(User_Score,Global_Sales)
```





```
# Create scatterplot matrix using ggplot() function
ggplot(Video.Game, aes(x = Year_of_Release, y = Global_Sales, color = Genre))
+
  geom_point() +
  labs(x = "Year of Release", y = "Global Sales", color = "Genre") +
  ggtitle("Scatterplot Matrix of Global Sales by Genre") +
  theme(plot.title = element_text(hjust = 0.5))
```

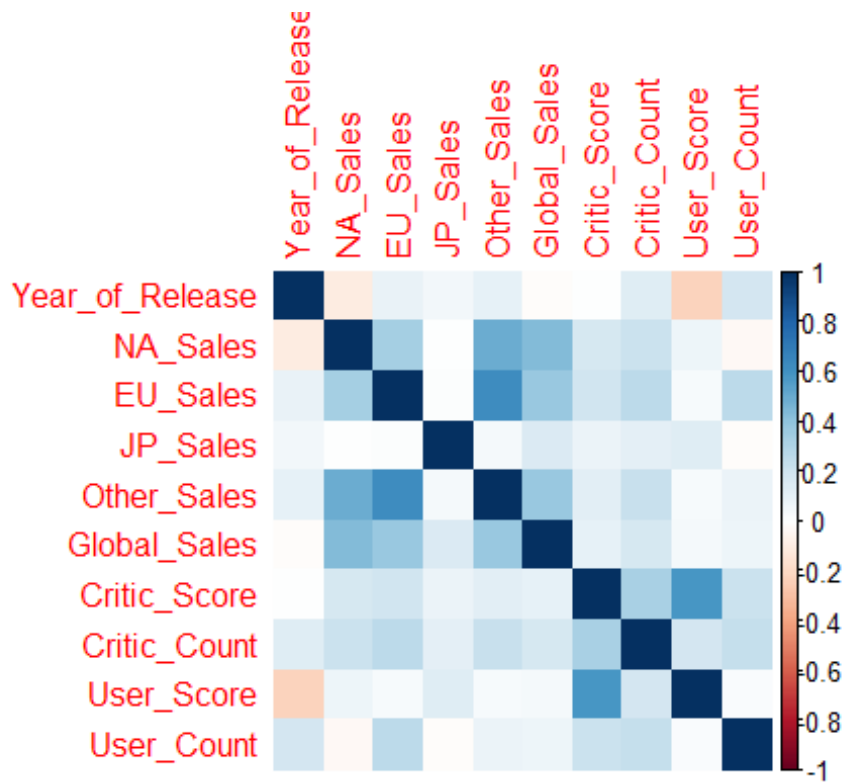
Scatterplot Matrix of Global Sales by Genre



```
# Generate boxplots for numeric variables
ggplot(Video.Game, aes(x = Genre, y = Global_Sales)) +
  geom_boxplot() +
  labs(x = "Genre", y = "Global Sales") +
  ggtitle("Boxplots of Global Sales by Genre") +
  theme(plot.title = element_text(hjust = 0.5)) +
  theme_bw()
```




```
# Pairwise correlation between all numeric variables:
corrplot(cor(Video.Game[, sapply(Video.Game, is.numeric)], use =
"complete.obs"), method = "color")
```



```
# Create scatterplot matrix using ggplot() function
ggplot(Video.Game, aes(x = Critic_Score, y = Global_Sales, color = Rating)) +
  geom_point() +
  ggtitle("Relationship Between Critic Score and Global Sales, by Rating") +
  xlab("Critic Score") +
  ylab("Global Sales")
```

Relationship Between Critic Score and Global Sales, by



```
stargazer(Video.Game, title = "Summary Statistics",
          type = "text", summary.stat = c("mean", "sd", "min", "max"))
```

```
##
## Summary Statistics
## =====
## Statistic      Mean      St. Dev.  Min    Max
## -----
## Year_of_Release 2,007.553  4.111    2,000  2,016
## NA_Sales        0.156    0.155    0.000  0.940
## EU_Sales        0.079    0.101    0.000  0.930
## JP_Sales        0.022    0.068    0.000  0.740
## Other_Sales     0.027    0.036    0.000  0.540
## Global_Sales    0.981    0.469   -0.575  2.568
## Critic_Score    67.962   13.571    13     96
## Critic_Count    25.574   16.220     3    106
## User_Score      7.082    1.468    0.500  9.600
## User_Count     102.719  388.330     4   10,665
## -----
```

Simple Linear Regression - User Score

```
# Fit a simple linear regression model
lm.fit <- lm(Global_Sales ~ User_Score, data = Video.Game)

# Get summary of the model
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Global_Sales ~ User_Score, data = Video.Game)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5780 -0.3167 -0.0091  0.3170  1.5938
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.871663   0.031109  28.020 < 2e-16 ***
## User_Score   0.015481   0.004301   3.599 0.000322 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4687 on 5508 degrees of freedom
## Multiple R-squared:  0.002346,    Adjusted R-squared:  0.002165
## F-statistic: 12.95 on 1 and 5508 DF,  p-value: 0.0003223

# Create a sequence of User_Score values for prediction
x_seq <- seq(from = min(Video.Game$User_Score), to =
max(Video.Game$User_Score), by = 1)

# Get the predicted values from the model
y_hat <- predict(lm.fit, newdata = list(User_Score = x_seq))


# Calculate the residuals
residuals <- Video.Game$Global_Sales - y_hat

# Calculate the RMSE
RMSE <- sqrt(mean(residuals^2))

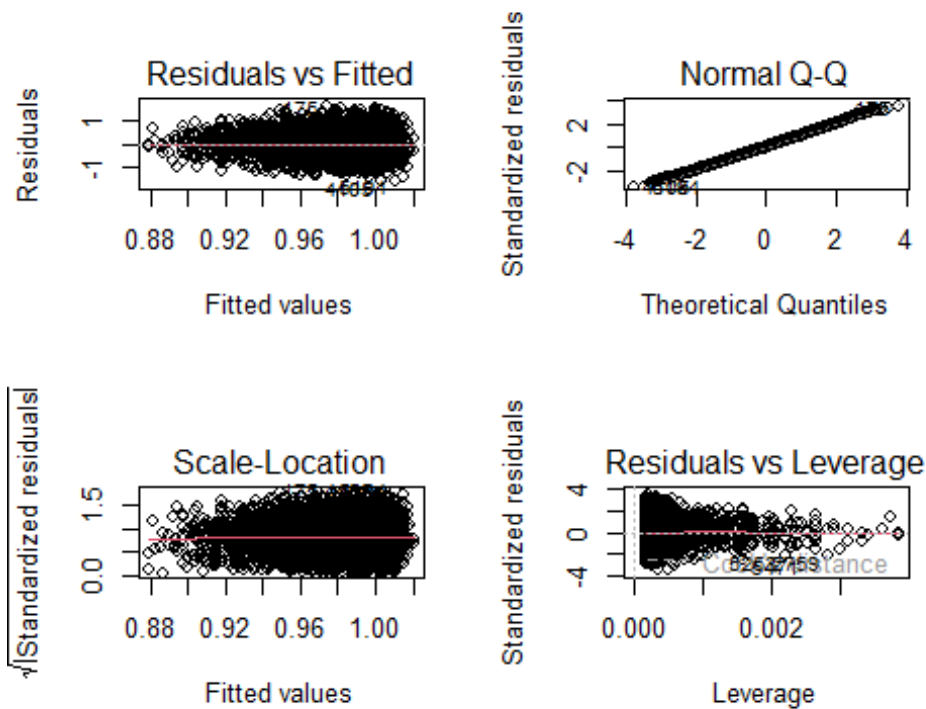
# Print the RMSE
cat("The RMSE of User Score for the simple regression model is:", RMSE, "\n")
## The RMSE of User Score for the simple regression model is: 0.4720381

# We will now plot Global_Sales and User_Score along with the Least squares
regression line
plot(User_Score, Global_Sales, pch = 20, col = "blue", main = "Video Game
Sales vs. User Score")

# Plot the simple regression line
lines(x_seq, y_hat, col = "red", lwd = 2)
```



```
# Produce four diagnostic plots  
par(mfrow = c(2,2))  
plot(lm.fit)
```



Simple Linear Regression - Critic Score

```
# Fit a simple linear regression model
lm.fit <- lm(Global_Sales ~ Critic_Score, data = Video.Game)

# Get summary of the model
summary(lm.fit)
##
## Call:
## lm(formula = Global_Sales ~ Critic_Score, data = Video.Game)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.61045 -0.31863 -0.00187  0.31152  1.58623
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.7390682  0.0321127  23.015  < 2e-16 ***
## Critic_Score 0.0035641  0.0004634   7.692 1.71e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4667 on 5508 degrees of freedom
## Multiple R-squared:  0.01063,    Adjusted R-squared:  0.01045
## F-statistic: 59.16 on 1 and 5508 DF,  p-value: 1.708e-14
```



```

# Create a sequence of Critic Score values for prediction
x_seq <- seq(from = min(Video.Game$Critic_Score), to =
max(Video.Game$Critic_Score), by = 1)

# Get the predicted values from the model
y_hat <- predict(lm.fit, newdata = list(Critic_Score = x_seq))

# Calculate the residuals
residuals <- Video.Game$Global_Sales - y_hat
## Warning in Video.Game$Global_Sales - y_hat: Longer object length is not a
multiple of shorter object length

# Calculate the RMSE
RMSE <- sqrt(mean(residuals^2))

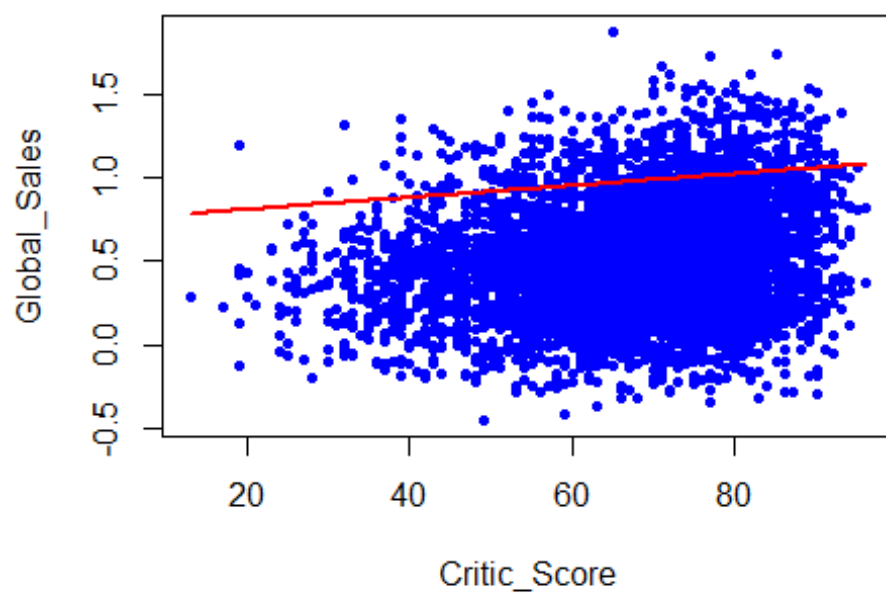
# Print the RMSE
cat("The RMSE of Critic Score for the simple regression model is:", RMSE,
"\n")
## The RMSE of Critic Score for the simple regression model is: 0.4779891

# We will now plot Global_Sales and Critic_Score along with the Least squares
regression line
plot(Critic_Score, Global_Sales, pch = 20, col = "blue", main = "Video Game
Sales vs. Critic Score")

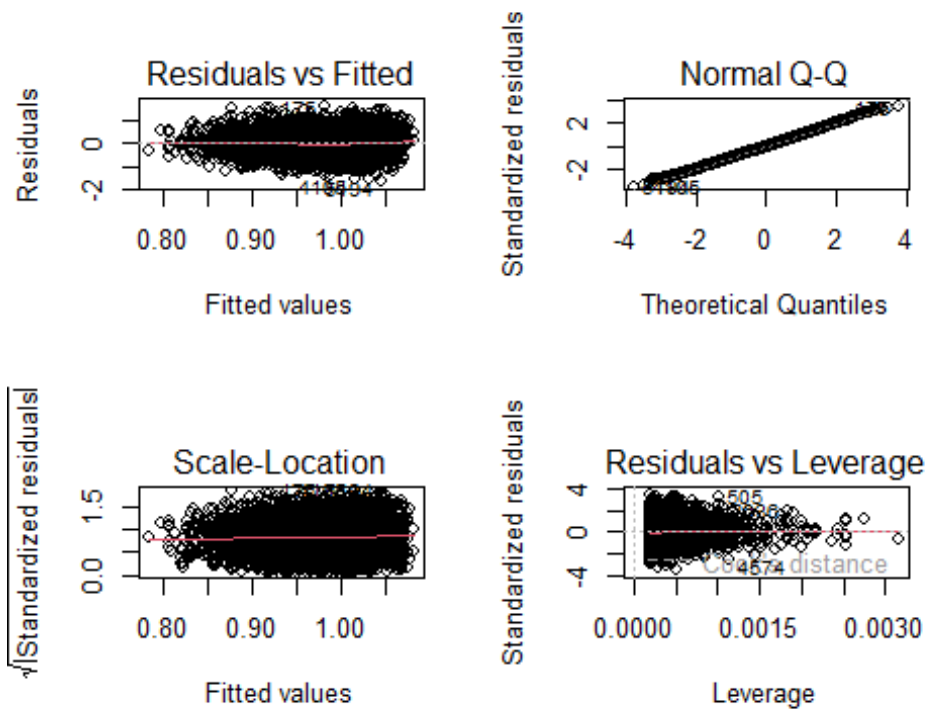
# Plot the simple regression line
lines(x_seq, y_hat, col = "red", lwd = 2)

```

Video Game Sales vs. Critic Score



```
# Produce four diagnostic plots  
par(mfrow = c(2,2))  
plot(lm.fit)
```



Simple Linear Regression - Confidence Intervals and Prediction Intervals - Critic_Score

```
# Indicating the dataset for the linear regression
lm.fit <- lm(Global_Sales ~ Critic_Score, data = Video.Game)

# Create a sequence of Critic Score values that are present in the dataset
critic_score_seq <- seq(from = min(Video.Game$Critic_Score), to =
max(Video.Game$Critic_Score), by = 1)

# Predict the Global Sales for each value in the sequence with 95% confidence
interval
confid <- predict(lm.fit, newdata = data.frame(Critic_Score =
critic_score_seq), interval = "confidence", level = 0.95)

# View the Confidence Interval for Global Sales
print(head(confid))
##           fit          lwr          upr
## 1 0.7854013 0.7339759 0.8368266
## 2 0.7889654 0.7384214 0.8395093
## 3 0.7925294 0.7428659 0.8421929
## 4 0.7960935 0.7473094 0.8448776
## 5 0.7996576 0.7517519 0.8475633
## 6 0.8032217 0.7561932 0.8502502
```

```

# Predict the Global Sales for each value in the sequence with prediction interval
pred <- predict(lm.fit, newdata = data.frame(Critic_Score =
critic_score_seq), interval = "prediction")

# View the Prediction Interval for Global Sales
print(head(pred))
##           fit           lwr           upr
## 1 0.7854013 -0.1310078 1.701810
## 2 0.7889654 -0.1273947 1.705325
## 3 0.7925294 -0.1237825 1.708841
## 4 0.7960935 -0.1201711 1.712358
## 5 0.7996576 -0.1165607 1.715876
## 6 0.8032217 -0.1129512 1.719395

```

Simple Linear Regression - Log transformation - Critic Score

```

# Log transformation
Video.Game$log_Critic_Score <- log(Video.Game$Critic_Score)

# Fit a simple linear regression model with Log-transformed Critic Score
lm.fit_log <- lm(Global_Sales ~ log_Critic_Score, data = Video.Game)

# Get summary of the model
summary(lm.fit_log)
##
## Call:
## lm(formula = Global_Sales ~ log_Critic_Score, data = Video.Game)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.59616 -0.31848 -0.00317  0.31241  1.58163
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.17281    0.11413   1.514    0.13
## log_Critic_Score 0.19273    0.02716   7.095 1.46e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4671 on 5508 degrees of freedom
## Multiple R-squared:  0.009056, Adjusted R-squared:  0.008876
## F-statistic: 50.34 on 1 and 5508 DF, p-value: 1.459e-12

# Create a sequence of Critic Score values for prediction
x_seq <- seq(from = min(Video.Game$log_Critic_Score), to =
max(Video.Game$log_Critic_Score), by = 1)

# Predict Global Sales for each value in the sequence with the Log model

```

```

y_hat_log <- predict(lm.fit_log, newdata = data.frame(log_Critic_Score =
x_seq))

# Calculate the residuals
residuals_log <- Video.Game$Global_Sales - predict(lm.fit_log)


# Calculate the RMSE
RMSE_log <- sqrt(mean(residuals_log^2))

# Print the RMSE
cat("The RMSE of the log model is:", RMSE_log)
## The RMSE of the Log model is: 0.4670106

# We will now plot Global_Sales and Log(Critic_Score) along with the Least
squares regression line
plot(Video.Game$log_Critic_Score, Video.Game$Global_Sales, pch = 20, col =
"blue", main = "Video Game Sales vs. Log(Critic Score)")

# Plot the log regression line
lines(x_seq, y_hat_log, col = "red", lwd = 2)

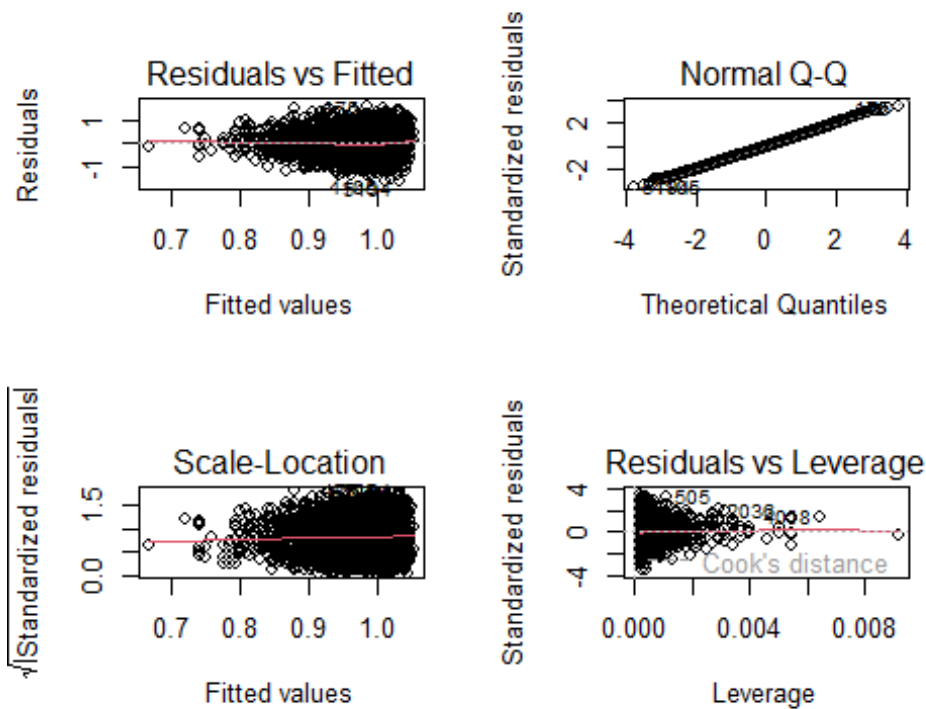
```



```

# Produce four diagnostic plots
par(mfrow = c(2,2))
plot(lm.fit_log)

```



Simple Linear Regression - Log transformation - User Score

Log transformation

```
Video.Game$log_User_Score <- log(Video.Game$User_Score)
```

Fit a simple linear regression model with Log-transformed Critic Score

```
lm.fit_log <- lm(Global_Sales ~ log_User_Score, data = Video.Game)
```

Get summary of the model

```
summary(lm.fit_log)
```

```
##
```

```
## Call:
```

```
## lm(formula = Global_Sales ~ log_User_Score, data = Video.Game)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -1.57363 -0.31682 -0.00997  0.31728  1.58975
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)    0.82175    0.04568  17.989 < 2e-16 ***
```

```
## log_User_Score  0.08275    0.02347   3.526 0.000425 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 0.4687 on 5508 degrees of freedom
```

```

## Multiple R-squared:  0.002253,   Adjusted R-squared:  0.002071
## F-statistic: 12.44 on 1 and 5508 DF,  p-value: 0.0004247

# Create a sequence of Critic Score values for prediction
x_seq <- seq(from = min(Video.Game$log_User_Score), to =
max(Video.Game$log_User_Score), by = 1)

# Predict Global Sales for each value in the sequence with the Log model
y_hat_log <- predict(lm.fit_log, newdata = data.frame(log_User_Score =
x_seq))

# Calculate the residuals
residuals_log <- Video.Game$Global_Sales - predict(lm.fit_log)

# Calculate the RMSE
RMSE_log <- sqrt(mean(residuals_log^2))


# Print the RMSE
cat("The RMSE of the log model is:", RMSE_log)
## The RMSE of the Log model is: 0.468611

# We will now plot Global_Sales and Log(Critic_Score) along with the Least
squares regression line
plot(Video.Game$log_User_Score, Video.Game$Global_Sales, pch = 20, col =
"blue", main = "Video Game Sales vs. Log(User Score)")

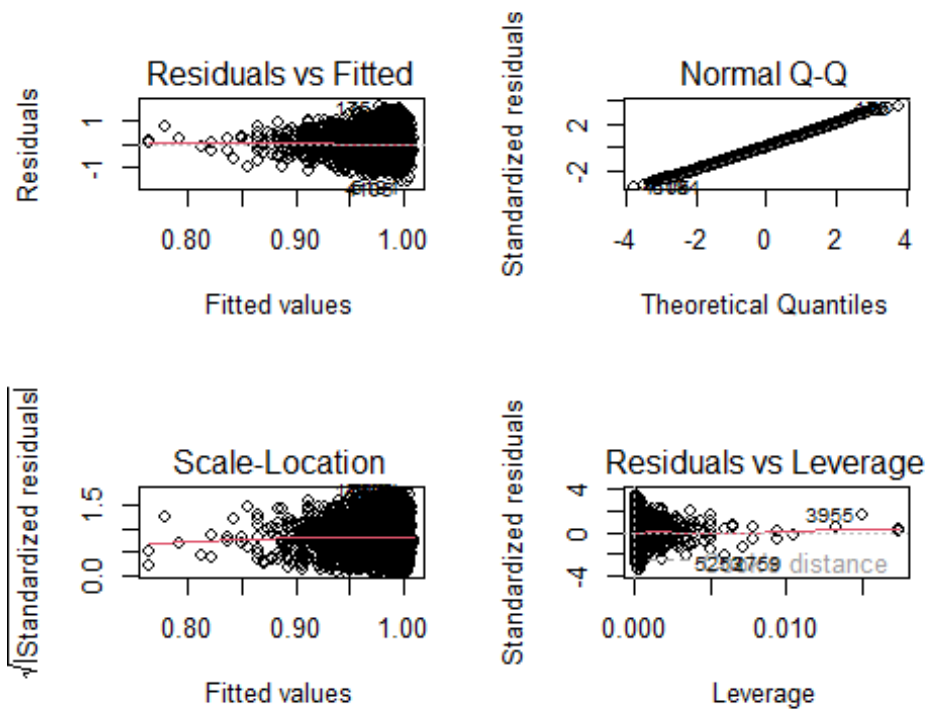
# Plot the Log regression line
lines(x_seq, y_hat_log, col = "red", lwd = 2)

```

Video Game Sales vs. Log(User Score)



```
# Produce four diagnostic plots  
par(mfrow = c(2,2))  
plot(lm.fit_log)
```

Logistic Regression - Genre

```
set.seed(123)
```

```
# Create binary variable based on Global_Sales threshold of revised_global_sales_mean
```

```
Video.Game$hit_median <- ifelse(Video.Game$Global_Sales >= revised_global_sales_mean, 1, 0)
```

```
# fit the logistic regression model
```

```
glm.fits <- glm(hit_median ~ Genre, family = binomial, data = Video.Game)
```

```
# Get summary of the model
```

```
summary(glm.fits)
```

```
##
```

```
## Call:
```

```
## glm(formula = hit_median ~ Genre, family = binomial, data = Video.Game)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -2.0849   0.4915   0.5562   0.5888   0.6681
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    1.72347    0.07638  22.564  <2e-16 ***
## GenreAdventure  0.17208    0.20440   0.842   0.3998
```

```

## GenreFighting      -0.08984      0.17386     -0.517      0.6053
## GenreMisc           0.21447      0.19394      1.106      0.2688
## GenrePlatform       0.06449      0.17990      0.358      0.7200
## GenrePuzzle         -0.33718      0.26763     -1.260      0.2077
## GenreRacing         -0.15486      0.14470     -1.070      0.2845
## GenreRole-Playing  -0.05905      0.13697     -0.431      0.6664
## GenreShooter        0.07543      0.13552      0.557      0.5778
## GenreSimulation     -0.06892      0.19073     -0.361      0.7179
## GenreSports         0.32914      0.13823      2.381      0.0173 *
## GenreStrategy      -0.30762      0.17518     -1.756      0.0791 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 4639.1  on 5509  degrees of freedom
## Residual deviance: 4620.0  on 5498  degrees of freedom
## AIC: 4644
##
## Number of Fisher Scoring iterations: 4

# Create a sequence of Genre values for prediction
x_seq <- levels(Video.Game$Genre)

# Get the predicted values from the model
y_hat <- predict(glm.fits, newdata = list(Genre = x_seq), type = "response")

# Calculate the residuals
residuals <- Video.Game$Global_Sales - y_hat
## Warning in Video.Game$Global_Sales - y_hat: Longer object length is not a
multiple of shorter object length


# Calculate the RMSE
RMSE <- sqrt(mean(residuals^2))

# Print the RMSE
cat("The RMSE of Genre for the logistic regression model is:", RMSE, "\n")
## The RMSE of Genre for the logistic regression model is: 0.488807

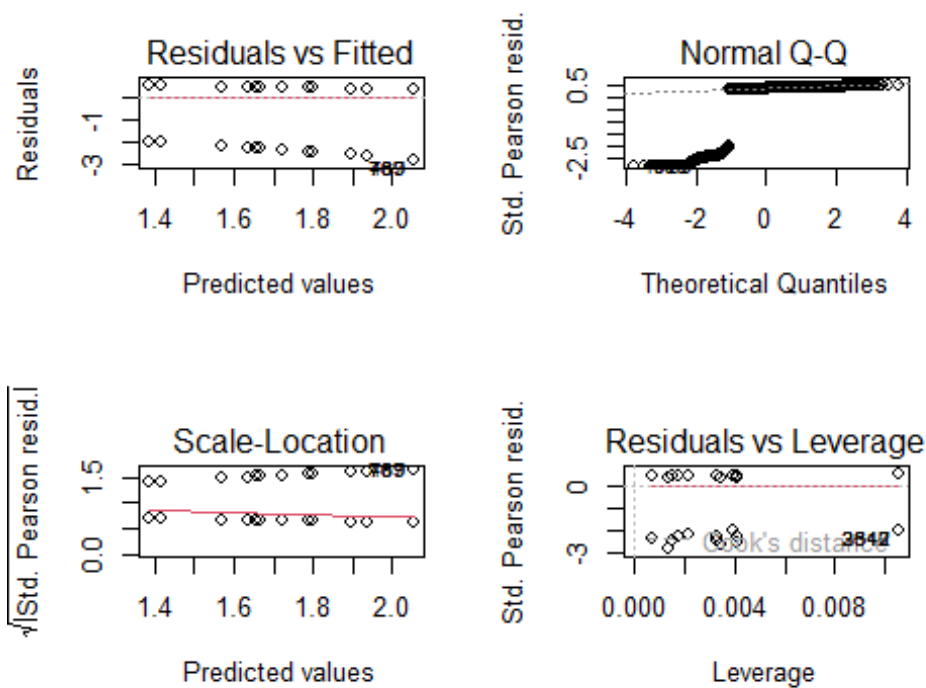
# We will now plot Global_Sales and Genre along with the Logistic regression
curve
plot(Genre, Global_Sales, pch = 20, col = "blue", main = "Video Game Sales
vs. Genre")

# Plot the logistic regression curve
lines(x_seq, y_hat, col = "red", lwd = 2)
## Warning in xy.coords(x, y): NAs introduced by coercion

```



```
# Produce four diagnostic plots  
par(mfrow = c(2,2))  
plot(glm.fits)
```



Logistic Regression - Platform

```
set.seed(123)
```

```
# Create binary variable based on Global_Sales threshold of
revised_global_sales_mean
```

```
Video.Game$hit_median <- ifelse(Video.Game$Global_Sales >=
revised_global_sales_mean, 1, 0)
```

```
# fit the logistic regression model
```

```
glm.fits <- glm(hit_median ~ Platform, family = binomial, data = Video.Game)
```

```
# Get summary of the model
```

```
summary(glm.fits)
```

```
##
```

```
## Call:
```

```
## glm(formula = hit_median ~ Platform, family = binomial, data = Video.Game)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -2.2618  0.4890  0.5152  0.6191  0.8446
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   1.91172    0.25984   7.357 1.88e-13 ***
## PlatformDC   -1.06442    0.73737  -1.444  0.14887
```

```

## PlatformDS    -0.30541    0.29373   -1.040    0.29845
## PlatformGBA   -0.14322    0.32846   -0.436    0.66282
## PlatformGC    -0.17968    0.30513   -0.589    0.55596
## PlatformPC    -0.75521    0.27609   -2.735    0.00623 **
## PlatformPS     0.15197    0.45650    0.333    0.73920
## PlatformPS2    0.04072    0.27928    0.146    0.88407
## PlatformPS3    0.56541    0.30451    1.857    0.06335 .
## PlatformPS4   -0.41779    0.32346   -1.292    0.19648
## PlatformPSP   -0.40406    0.29533   -1.368    0.17126
## PlatformPSV   -0.43445    0.35498   -1.224    0.22099
## PlatformWii    0.09852    0.30642    0.322    0.74781
## PlatformWiiU  -0.19874    0.41812   -0.475    0.63456
## PlatformX360   0.07776    0.28685    0.271    0.78634
## PlatformXB    -0.35709    0.28457   -1.255    0.20953
## PlatformXOne   0.10318    0.38532    0.268    0.78886
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 4639.1  on 5509  degrees of freedom
## Residual deviance: 4553.0  on 5493  degrees of freedom
## AIC: 4587
##
## Number of Fisher Scoring iterations: 5

# Create a sequence of Platform values for prediction
x_seq <- levels(Video.Game$Platform)

# Get the predicted values from the model
y_hat <- predict(glm.fits, newdata = list(Platform = x_seq), type =
"response")

# Calculate the residuals
residuals <- Video.Game$Global_Sales - y_hat
## Warning in Video.Game$Global_Sales - y_hat: Longer object length is not a
multiple of shorter object length


# Calculate the RMSE
RMSE <- sqrt(mean(residuals^2))

# Print the RMSE
cat("The RMSE of Platform for the logistic regression model is:", RMSE, "\n")
## The RMSE of Platform for the Logistic regression model is: 0.491246

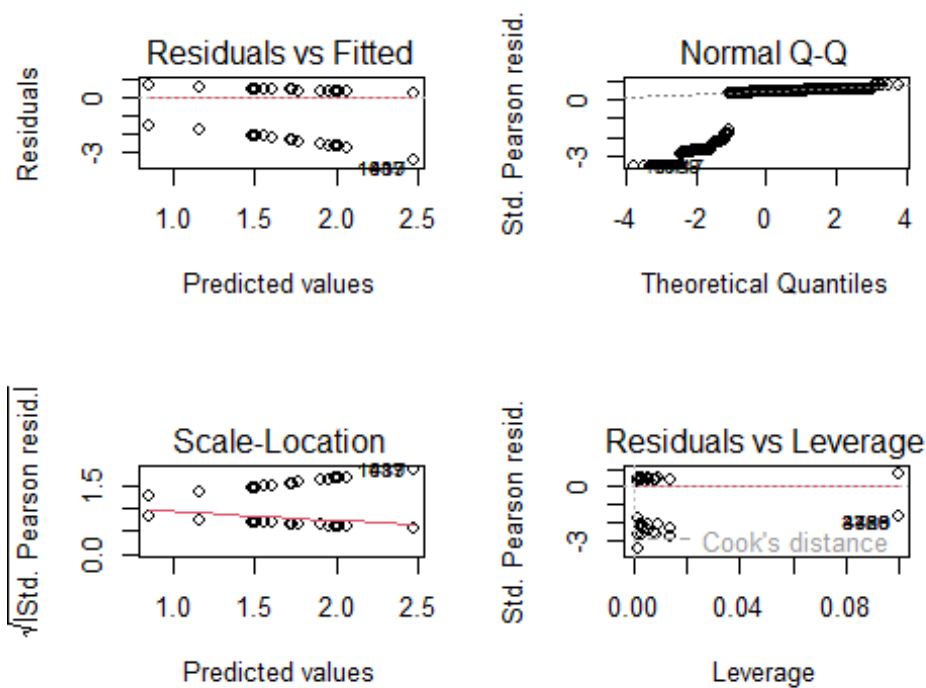
# We will now plot Global_Sales and Platform along with the Logistic
regression curve
plot(Platform, Global_Sales, pch = 20, col = "blue", main = "Video Game Sales
vs. Platform")

```

```
# Plot the logistic regression curve
lines(x_seq, y_hat, col = "red", lwd = 2)
## Warning in xy.coords(x, y): NAs introduced by coercion
```



```
# Produce four diagnostic plots
par(mfrow = c(2,2))
plot(glm.fits)
```



Multiple Regression - NA_Sales, Genre, Critic_Score, & Critic_Count

```
set.seed(123)
```

```
# fit the linear regression model
```

```
lm.fits <- lm(Global_Sales ~ NA_Sales + Genre + Critic_Score + Critic_Count,
data = Video.Game)
```

```
# Get summary of the model
```

```
summary(lm.fits)
```

```
##
```

```
## Call:
```

```
## lm(formula = Global_Sales ~ NA_Sales + Genre + Critic_Score +
##     Critic_Count, data = Video.Game)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -1.46876 -0.28017  0.00244  0.28316  1.49236
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.7226853   0.0300009   24.089   < 2e-16 ***
## NA_Sales       1.2739096   0.0389434   32.712   < 2e-16 ***
## GenreAdventure  0.0468104   0.0293363    1.596   0.1106
## GenreFighting  -0.0598050   0.0269160   -2.222   0.0263 *
## GenreMisc      -0.0427569   0.0275468   -1.552   0.1207
```

```
## GenrePlatform      -0.0285683  0.0267512 -1.068  0.2856
## GenrePuzzle        -0.0561119  0.0447225 -1.255  0.2097
## GenreRacing        -0.0211052  0.0227701 -0.927  0.3540
## GenreRole-Playing  0.0086815  0.0211274  0.411  0.6812
## GenreShooter       -0.0049118  0.0202135 -0.243  0.8080
## GenreSimulation    -0.0295905  0.0294779 -1.004  0.3155
## GenreSports        -0.0155668  0.0198832 -0.783  0.4337
## GenreStrategy      -0.0280031  0.0290330 -0.965  0.3348
## Critic_Score       0.0001778  0.0004554  0.391  0.6962
## Critic_Count       0.0023434  0.0003919  5.980 2.37e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4206 on 5495 degrees of freedom
## Multiple R-squared:  0.1984, Adjusted R-squared:  0.1964
## F-statistic: 97.17 on 14 and 5495 DF,  p-value: < 2.2e-16

# Create a sequence of Genre, Critic_Score, Critic_Count and NA_Sales values
for prediction
subset_data <- na.omit(Video.Game[c("Genre", "Critic_Score", "Critic_Count",
"NA_Sales", "Global_Sales")])

x_seq <- expand.grid(
  Genre = unique(subset_data$Genre),
  Critic_Score = seq(min(subset_data$Critic_Score),
max(subset_data$Critic_Score), length.out = 50),
  Critic_Count = seq(min(subset_data$Critic_Count),
max(subset_data$Critic_Count), length.out = 50),
  NA_Sales = seq(min(subset_data$NA_Sales), max(subset_data$NA_Sales),
length.out = 50)
)

# Get the predicted values from the model
y_hat <- predict(lm.fits, newdata = x_seq)

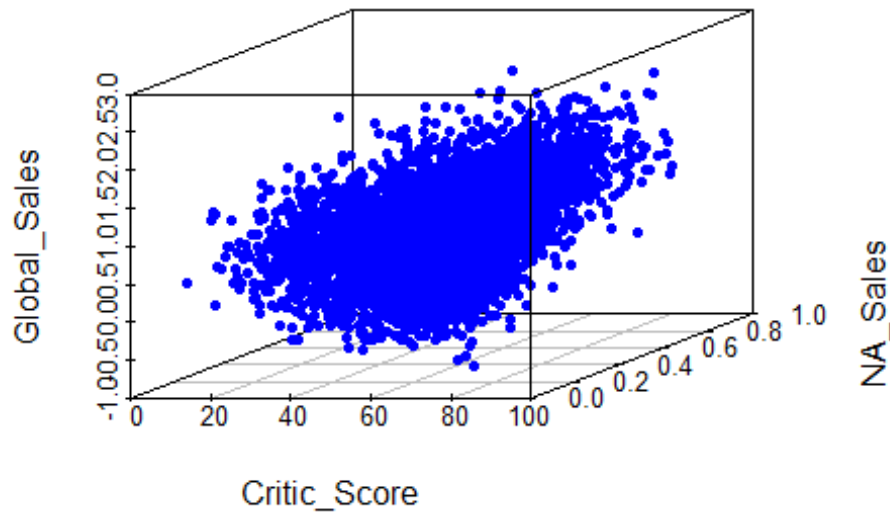
# Calculate the residuals
residuals <- subset_data$Global_Sales - y_hat
## Warning in subset_data$Global_Sales - y_hat: Longer object length is not a
multiple of shorter object length

# Calculate the RMSE
RMSE <- sqrt(mean(residuals^2))

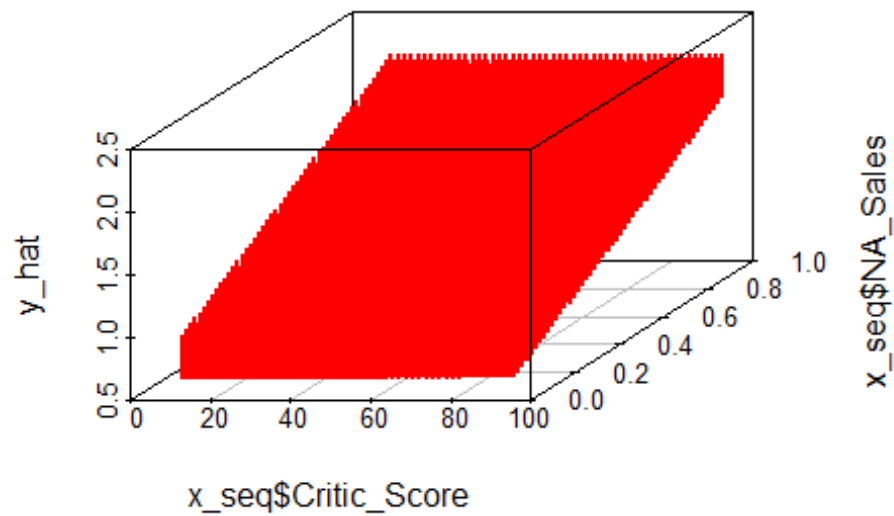
# Print the RMSE
cat("The RMSE of Genre, Critic_Score, Critic_Count, and NA_Sales for the
linear regression model is:", RMSE, "\n")
## The RMSE of Genre, Critic_Score, Critic_Count, and NA_Sales for the Linear
regression model is: 0.7486262
```




```
# Create a 3D scatter plot with Critic_Score, NA_Sales, and Global_Sales
scatterplot3d(subset_data$Critic_Score, subset_data$NA_Sales,
subset_data$Global_Sales,
              color = "blue", pch = 20, xlab = "Critic_Score", ylab =
"NA_Sales", zlab = "Global_Sales")
```



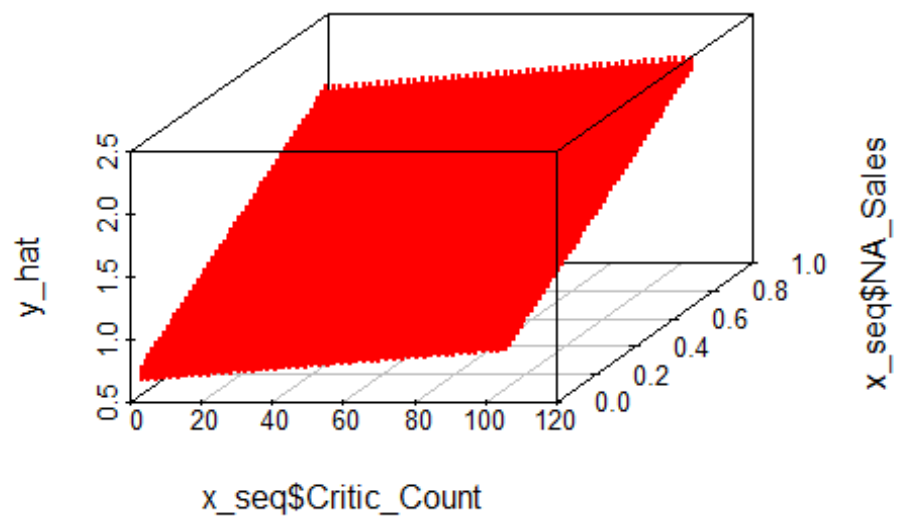
```
# Add linear regression curve to the 3D scatter plot
scatterplot3d(x_seq$Critic_Score, x_seq$NA_Sales, y_hat,
              color = "red", add = TRUE, type = "l", lwd = 2)
## Warning in title(main, sub, ...): "add" is not a graphical parameter
## Warning in plot.xy(xy.coords(x, y), type = type, ...): "add" is not a
graphical parameter
```







```
# Create a 3D scatter plot with Critic_Count, NA_Sales, and Global_Sales
scatterplot3d(subset_data$Critic_Count, subset_data$NA_Sales,
subset_data$Global_Sales,
              color = "blue", pch = 20, xlab = "Critic_Count", ylab =
"NA_Sales", zlab = "Global_Sales")
```



```
# Add linear regression curve to the 3D scatter plot
scatterplot3d(x_seq$Critic_Count, x_seq$NA_Sales, y_hat,
              color = "red", add = TRUE, type = "l", lwd = 2)
## Warning in title(main, sub, ...): "add" is not a graphical parameter
## Warning in title(main, sub, ...): "add" is not a graphical parameter
```



```
# Produce diagnostic plots  
par(mfrow = c(2,2))  
plot(lm.fits)
```



Multiple Regression - Regional_Sales, Genre, Critic_Score

```
set.seed(123)
```

```
# Create weights based on each region's contribution to sales (excluding Other_Region_Sales)
```

```
weights <- (Video.Game$NA_Sales + Video.Game$EU_Sales + Video.Game$JP_Sales) / sum(Video.Game$NA_Sales + Video.Game$EU_Sales + Video.Game$JP_Sales)
```

```
# Replace missing or negative weights with 0
```

```
weights[is.na(weights) | weights < 0] <- 0
```

```
# Fit the linear regression model with weights
```

```
lm.fits <- lm(Global_Sales ~ Critic_Score + Genre + NA_Sales + EU_Sales + JP_Sales, data = Video.Game, weights = weights)
```

```
# Get summary of the model
```

```
summary(lm.fits)
```

```
##
```

```
## Call:
```

```
## lm(formula = Global_Sales ~ Critic_Score + Genre + NA_Sales + EU_Sales + JP_Sales, data = Video.Game, weights = weights)
```

```
##
```

```
## Weighted Residuals:
```

```
##      Min      1Q    Median      3Q      Max
```

```
## -0.0258025 -0.0028484 -0.0000271 0.0028017 0.0240309
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.7902813   0.0312229   25.311 < 2e-16 ***
## Critic_Score  -0.0013355   0.0004462   -2.993 0.00277 **
## GenreAdventure  0.0471782   0.0330955    1.426 0.15406
## GenreFighting  -0.0371957   0.0243018   -1.531 0.12593
## GenreMisc       0.0147476   0.0247751    0.595 0.55169
## GenrePlatform  -0.0124587   0.0246529   -0.505 0.61333
## GenrePuzzle     0.0460494   0.0488207    0.943 0.34560
## GenreRacing    -0.0030512   0.0217440   -0.140 0.88841
## GenreRole-Playing 0.0298905   0.0204237    1.464 0.14338
## GenreShooter    0.0330287   0.0190219    1.736 0.08256 .
## GenreSimulation -0.0451897   0.0274588   -1.646 0.09988 .
## GenreSports     0.0394924   0.0176845    2.233 0.02558 *
## GenreStrategy   0.0573543   0.0339483    1.689 0.09119 .
## NA_Sales        1.0322691   0.0300013   34.407 < 2e-16 ***
## EU_Sales        1.1885530   0.0418598   28.394 < 2e-16 ***
## JP_Sales        0.9964170   0.0569562   17.494 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.005324 on 5494 degrees of freedom
## Multiple R-squared:  0.2966, Adjusted R-squared:  0.2946
## F-statistic: 154.4 on 15 and 5494 DF,  p-value: < 2.2e-16

# Create a sequence of Genre, Critic_Score, NA_Sales, EU_Sales, JP_Sales
values for prediction
subset_data <- na.omit(Video.Game[c("Genre", "Critic_Score", "NA_Sales",
"EU_Sales", "JP_Sales", "Global_Sales")])
x_seq <- expand.grid(
  Genre = unique(subset_data$Genre),
  Critic_Score = seq(min(subset_data$Critic_Score),
max(subset_data$Critic_Score), length.out = 25),
  NA_Sales = seq(min(subset_data$NA_Sales), max(subset_data$NA_Sales),
length.out = 25),
  EU_Sales = seq(min(subset_data$EU_Sales), max(subset_data$EU_Sales),
length.out = 25),
  JP_Sales = seq(min(subset_data$JP_Sales), max(subset_data$JP_Sales),
length.out = 25)
)

# Get the predicted values from the model
y_hat <- predict(lm.fits, newdata = x_seq)

# Calculate the residuals
residuals <- subset_data$Global_Sales - y_hat
## Warning in subset_data$Global_Sales - y_hat: Longer object length is not a
multiple of shorter object length
```


```

# Calculate the RMSE
RMSE <- sqrt(mean(residuals^2))

# Print the RMSE
cat("The RMSE of Genre, Critic_Score, NA_Sales, EU_Sales, JP_Sales for the
linear regression model is:", RMSE, "\n")
## The RMSE of Genre, Critic_Score, NA_Sales, EU_Sales, JP_Sales for the
linear regression model is: 1.343167

# Create a 3D scatter plot with Genre, NA_Sales, and Global_Sales
scatterplot3d(subset_data$Genre, subset_data$NA_Sales,
subset_data$Global_Sales,
              color = "blue", pch = 20, xlab = "Genre", ylab = "NA_Sales",
zlab = "Global_Sales")


```




```

# Add linear regression curve to the 3D scatter plot
scatterplot3d(x_seq$Genre, x_seq$NA_Sales, y_hat,
              color = "red", add = TRUE, type = "l", lwd = 2)
## Warning in title(main, sub, ...): "add" is not a graphical parameter
## Warning in plot.xy(xy.coords(x, y), type = type, ...): "add" is not a
graphical parameter


```




```
# Create a 3D scatter plot with NA_Sales, EU_Sales, and Global_Sales
scatterplot3d(subset_data$NA_Sales, subset_data$EU_Sales,
subset_data$Global_Sales,
               color = "blue", pch = 20, xlab = "NA_Sales", ylab = "EU_Sales",
zlab = "Global_Sales")
```


```
# Add linear regression curve to the 3D scatter plot
scatterplot3d(x_seq$NA_Sales, x_seq$EU_Sales, y_hat,
              color = "red", add = TRUE, type = "l", lwd = 2)
## Warning in title(main, sub, ...): "add" is not a graphical parameter
## Warning in title(main, sub, ...): "add" is not a graphical parameter
```



```
# Create a 3D scatter plot with JP_Sales, EU_Sales, and Global_Sales
scatterplot3d(subset_data$JP_Sales, subset_data$EU_Sales,
subset_data$Global_Sales,
               color = "blue", pch = 20, xlab = "JP_Sales", ylab = "EU_Sales",
zlab = "Global_Sales")
```



```
# Add linear regression curve to the 3D scatter plot
scatterplot3d(x_seq$JP_Sales, x_seq$EU_Sales, y_hat,
              color = "red", add = TRUE, type = "l", lwd = 2)
## Warning in title(main, sub, ...): "add" is not a graphical parameter
## Warning in title(main, sub, ...): "add" is not a graphical parameter
```



```
# Produce diagnostic plots  
par(mfrow = c(2,2))  
plot(lm.fits)
```

